

CS3505 Final Project, Fall 2020

Team Exoplanets

Edward Barrowes, Corbin Gurnee, Ted Goodell,
Marley Stark, Ivan Burkett, Garrett Keefe

December 9, 2020

1 Overview

Our group project was to develop a web-scraping and data synthesizing tool that could use FFMPEG to create a scaled timeline of data from multiple sources. The project needed to be easily scaled, flexible, run remotely, and require minimal user input to generate timelines for arbitrary time windows. Additionally, we were tasked with making use of cloud storage to the best of our ability to support the architecture of our project.

We made use of multiple technologies in developing our tool, including:

- Amazon Web Services
- Docker
- GitHub
- Python
- Selenium
- FFMPEG
- C++

We were successful in meeting the aims of this project, including the educational aim of practicing AGILE development in a team based environment, all of which is outlined within this document.

2 Architecture

2.1 Organization

Scalability is an important goal for this project, which factors in to the organizational choices our team made. At the top level, all of our code is run on remote environments, and at lower levels is compartmentalized to allow for easy scalability. This structure is illustrated in figure 1

The main operations of the program are run within a virtual machine instance using a service called Elastic Compute Cloud (EC2), which is itself a part of the greater Amazon Web Service (AWS). Within the EC2 instance are two Docker images: one for data collection, and one for data synthesis.

The data collection image uses a web-scraping tool called Selenium, to collect images and data from internet sources. Selenium supports multiple languages, but we chose to use python to interface with its scripts. Excepting the data collected from twitch, all our data, both text based¹ and image based, was collected using Selenium. For twitch, we wrote C++ code which directly interfaces with the Twitch API using a built-in C++ library called `lcurl`. These data are stored as .png images in cloud storage, again provided by Amazon, called the Elastic File System (EFS).

Data synthesis occurs within a Docker image containing code we constructed, `framebuilder.cpp`. `framebuilder` assembles individual images collected in the EFS, combining them into 600 time-consecutive frames using FFMPEG. Subsequently, the frames are assembled into a minute long video (again using FFMPEG) showing all of the data.




Figure 1: UML Class diagram depicting project structure

2.2 Execution

Execution occurs on both a continuous and instantaneous basis. Data collection occurs autonomously, and continuously, requiring no direct user input. Data synthesis occurs only when queried by the user interacting with the EC2 instance, illustrated in figure 2.

¹In-between data collection and synthesis is a bit of formatting. Specifically, text based data (i.e. number of viewers on twitch) is converted to an .png file using a python script which interfaces directly with FFMPEG.

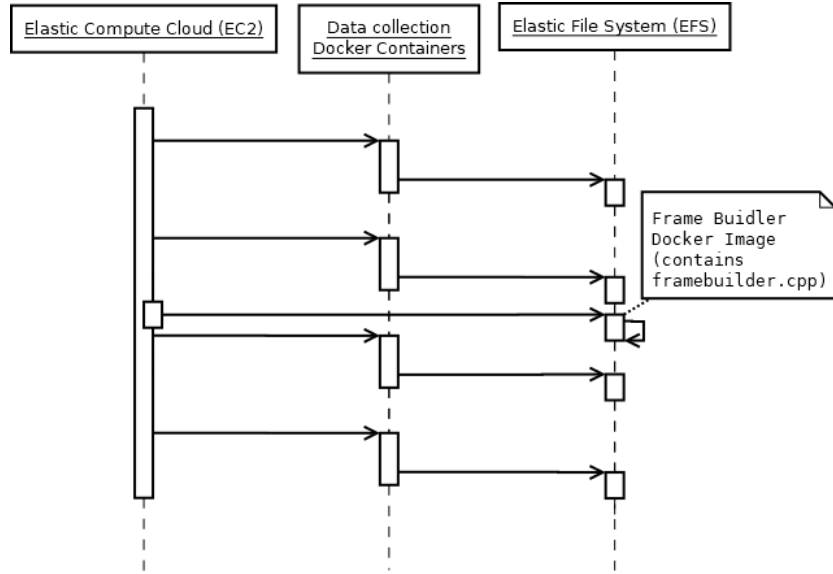


Figure 2: UML Sequence diagram depicting project execution

Autonomous data collection occurs at regular intervals, using the Linux command 'cron'. When data collection occurs, six Docker containers are instantiated from the same image, each containing python code which collects data from various sources. The image itself contains all the code that all containers could need, however each data source updates at a different frequency. To compensate for this, we instantiate containers at appropriate time intervals in order to collect data.

A user can query the EC2 instance to produce a video compilation of collected data by running a bash script within the virtual machine. Running the script launches a Docker container, which contains our framebuilder.cpp program, which assembles all the data together into a minute long video. The container itself has FFmpeg installed, since framebuilder.cpp relies on its use.

3 Development

3.1 Organizational Tools

To organize our project, we made use of GitHub and its associated functionality. All our code was backed by a git repository², and our development progress was tracked using GitHub Projects and Issue Tracker seen in figures 3 and 4.

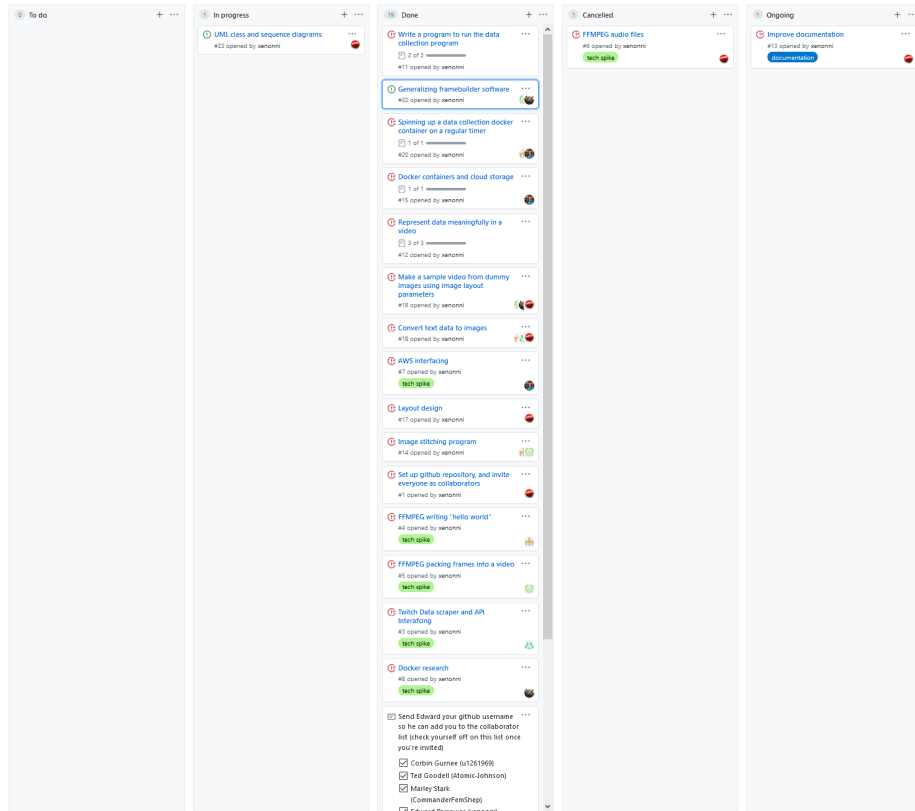


Figure 3: A screenshot showing the tracker, GitHub projects, which we used to organize development.

Additionally, we took regular project development meeting notes, which were also stored in the repository and have been attached in appendix A.

²https://github.com/xenonni/cs3505_space.data

to run off of the web scraper as well. After this, I worked with Ted to convert the data we have compiled from our web scrapers to images. Finally, I worked in limited capacities with most other members of the team to assist in any area they may need assistance with.

Ted Goodell - I worked with Ivan on using the ffmpeg libraries to combine multiple images into a single image. To do this, we wrote code that decoded several images from their files and put them into frame structs that we could then programmatically combine into a larger frame. We then used the ffmpeg libraries to encode that frame and put it into an output image. I also, wrote python scripts that grabbed text data output by the data grabbers, and used the ffmpeg command line utility to put that text into an image. The ffmpeg command used a built in filter called “drawtext” to draw the text on a black template image for that portion of the final output image.

Ivan Burkett - I was assigned to work on the frame builder program for this project. I worked with Ted to create a C++ program called image-smoosh.cpp that used ffmpeg to combine the 6 frame segments into a single frame, then saved that frame as a png image. After finishing image-smoosh.cpp, I refactored the program to make it cleaner and easier to read and saved it as image-smoosh2.cpp. After doing that, I used image-smoosh2.cpp as a base to create framebuilder.cpp with Garrett. As part of making framebuilder.cpp, I also created a wrapper class for the C++ class directory-iterator called peek-directory-iterator, which simplified usage for directory-iterator and added a peek functionality which was necessary for our file traversal algorithm. framebuilder.cpp uses peek-directory-iterators to scan through the data source folders for the one most up to date for the provided date during the frame building process. Due to the way AWS’s filesystem works, I ended up scrapping peek_directory_iterator. the final version loads all files in the directory and then sorts them, because AWS doesn’t have any useful ordering to the files and directory_iterator doesn’t ensure any kind of ordering. loading them all and sorting them was the only way to ensure the ordering.

Garrett Keefe - I first worked on gathering research on docker and how it works. Then I was assigned with Corbin to work on the data scrapers to collect data from our 6 sources. I wrote all of the .py files which used selenium to navigate and grab web resources while Corbin wrote the twitch.out file which uses C++ and the curl libraries to query twitch.tv for info from its API. After that I was assigned to working with Ivan to create the framebuilder program which utilize UNIX timestamps to assemble data from our data sources and compile them together into individual frames, allowing us to generate a list of frames which we use an ffmpeg command to assemble them into a video. I also helped Ted create all of the .sh files which acted as drivers for our data collection and converting to picture formats within the docker containers.

Edward Barrowes - Most of my work on this project was in documentation and organization. I took meeting notes whenever we met and discussed anything of significance. Additionally, I was responsible for setting up the GitHub repository and populating the issue tracker with tasks we had discussed during meetings (though individual team members were responsible for updating progress on tasks they had been assigned). In addition to providing logistical support to the team, I also designed the final layout for the video based on the requirements and constraints the various data sources had. Lastly, as we finished the project, I was responsible for compiling our design discussions and sketches (from weekly meetings) into more professional communication, including this document and the UML diagrams contained within.

Marley Stark - I worked on AWS to get our EC2, EFS, and lambda functions set up so we could have a virtual environment on the cloud to automate and run our entire program and data collection. I made an AWS Educate account so we could explore what AWS provides without having to worry about being charged, and it seemed to have access to everything we needed for the assignment. However the AWS Educate account does not let you schedule EC2 instances so we were unable to set up the Cloudwatch events needed to use lambda functions on a schedule. Our lambda functions work to turn on and off our EC2, we just couldn't schedule them to automate it. Taking advantage of Linux's cron daemon, I edited the crontab file in Cloud9 to schedule our data collection jobs, and it works fantastic! So overall I navigated the immense world of AWS and made a Cloud9 EC2 instance and mounted an EFS to it, learned about Lambda functions and Cloudwatch but was unable to schedule things because of the AWS Educate account restrictions, and instead learned about and implemented cron jobs to do our automated data collection tasks.

4 Results

According to the goals of our project, we have complete functionality. Our project works fully as intended because we can leave our EC2 instance running to collect data without having to even be signed into our AWS account. The timeline video is generated with minimal user input.

4.1 Running the Code

To make a video from the data you simply run the command
`./create_video.sh MM.DD.YYYY.HH.mm MM.DD.YYYY.HH.mm`
with the desired start and end times.



Figure 5: A screenshot from the completed timeline.

5 Appendix A: Meeting Notes

5.1 Meeting notes 10/29

- 1 Todo for next checkpoint (1 week from today)
- 2 -----
- 3
- 4 TA: Josh Nelson, Thursday 4:15 PM
- 5 Setup github repo
- 6 Using the tracking tool of choice (github projects) to setup a
 - backlog of work, broken into tasks that take 3-4 hours,
 - leading from start to finish. These can be updated and
 - revised throughout the entirety of the project.
- 7 Plan the first sprint
- 8 Assign tasks to people for the first sprint
- 9 Recommended: Spike tasks as part of the sprint
 - 10 1-2 hour long tasks
 - 11 example: draw a smiley face with ffmpeg as a .png
 - 12 The idea is so you know how to interface with a a part of the
 - system for future troubleshooting
- 13 Plan regularly scheduled standups. Everyone should be there,
 - everyone must share:

14 What they're working on
15 What they're stuck with
16 What they're doing next
17 Suggested one standup per sprint
18 Each person gets about 1-2 minutes for this
19 The timescale of the minute of output data is up to you
20 it could be as much as 48 hours compressed into a minute
21 another group using 1/day covid data was "pushing it", so
 ↪ probably try not to go lower frequency than that for more
 ↪ than one data source
22 Grabbing the most recent data point from a slow updating source
 ↪ would be fine
23 The most ideal would be to have data that updates on the scale
 ↪ of hours rather than days
24 Data seeding: ask the professor
25 We don't have to use the exact same time slot, but we should try
 ↪ and get the same TA

5.2 Meeting notes 11/3

1 Members Absent: None
2
3 Next meeting
4 TA: Josh Nelson, Thursday 4:15 PM
5
6 Tech spikes
7 Data scraper & API interfacing - Corbin
8 Continuous data monitoring
9 API Keys
10 ffmpeg
11 a program that writes the text "hello world" in a picture -
 ↪ Ted
12 take multiple frames and pack them into a video - Ivan
13 dealing with audio files - Edward
14 AWS or other cloud based computing - Marley
15 Somewhere, the docker container must run on a cloud service
 ↪ and use cloud storage to run the program
16 More docker research - Garrett
17 Running additional docker containers from a container, in
 ↪ order to run additional tasks?
18
19 Overall Project steps
20 Program to collect data from the internet in the short term
21 Identify data sources (complete)
22 Identify ways to get the data
23 scraper or API

24 Unified format to save data in
 25 might depend on ffmpeg
 26 Program to run the data collection program
 27 Involves deployment of docker containers
 28 Needs to run for a long time (about 24h) calling data
 ↪ collection program at regular intervals
 29 Store the data
 30 Persistent and accessible cloud storage where all the data is
 ↪ collected
 31 Represent data meaningfully in a video
 32 Visualize data
 33 Scale the time appropriately for a minute long video
 34 Output video file
 35
 36 Tasks
 37 Make github cards - Edward
 38 Contact TA about moving the meeting time to 3:30 PM, Thursdays
 ↪ - Marley
 39 Data sources
 40 Supernovae data -
 ↪ <http://www.rochesterastronomy.org/snimages/sndate.html>
 41 Updates once per day, with multiple entries (2-8)
 42 Does it update at a regular time?
 43 Might be too slow
 44 Weather data - <https://aprs.fi/weather/a/FW3937?range=day>
 45 API - figure out which of these works best
 46 <https://aprs.fi/page/api>
 47
 ↪ <https://weather.com/swagger-docs/ui/sun/v3/sunV3CurrentsOnDemand.json>
 48 Lightning data -
 ↪ https://ghrc.nsstc.nasa.gov/hydro/details/isslis_v1_nrt
 49 Twitch viewer data
 50 updates very rapidly
 51 API - <https://dev.twitch.tv/docs/api/>
 52 ISS location tracker -
 ↪ https://www.esa.int/Science_Exploration/Human_and_Robotic_Exploration/International
 53 updates very rapidly
 54 <https://api.nasa.gov/>
 55 Garrett will pick a data source from this API
 56
 57 Questions for TA
 58 Running additional docker containers from one container?
 59 Also, is this required? or can all the code run from a single
 ↪ docker container? (no)
 60 Any required documentation, beyond github projects?
 61 no additional documentation required

62 Do we need to pull actual images/audio from the web, or just
↳ make images/audio to represent data?
63 text to speech could count as pulling audio data from the
↳ internet

5.3 Meeting notes 11/10

1 Members absent: none
2
3 Tech spike updates
4
5 - Edward has done some research on ffmpeg audio, but isn't sure
↳ how much more we want to pursue it for our project in light
↳ of the most recent TA meeting
6 - Ivan had some demo code to show on packing frames into a
↳ video. There were some parts of the code that weren't well
↳ understood, so we all discussed together. Ivan also showed
↳ us a video that was generated by ffmpeg
7 - Marley looked into web services to host our project on, and
↳ applied for an AWS Educate Starter Account. We should get
↳ access in the next few days or so.
8 Garrett identified some features that we may potentially need
↳ from AWS:
9 - Elastic container storage
10 - Event handling
11 - Lambdas (maybe)
12 - Garrett reviewed some of the parameters for docker commands
13 --rm can be used to remove a container once it has run
14 Garrett also wrote a .bat file to run the docker container he was
↳ demonstrating, which was added to a separate branch of the
↳ repository
15 - Corbin looked through and knows what to do to implement API of
↳ twitch, will have it fully working by Thursday at the
↳ latest.
16 - Ted looked into using libavfilter to draw text on the screen.
↳ He is going to modify the video filter example on
17
18 Visual layout discussion
19 Have a circle which displays temperature as text in Salt Lake
↳ City, potentially scrape image from weather.com, center
↳ bottom
20 ISS image grab as well as latitude and longitude can go in the
↳ top left corner
21 Lightning strike map image can go in the top right corner
22 Supernova data on the bottom left
23 Twitch data on the bottom right

24 See meeting_notes_image_11_10.png for more detail




Figure 6: Image from meeting notes for 11/10

5.4 Meeting notes 11/12

- 1 One of the requirements is to write C++ code that interfaces with
 - FFMPEG at some point. Displaying weather data using a color
 - drawn by ffmpeg (orange for increase, blue for decrease)
 - could easily satisfy this requirement.
- 2 We discussed turning all data into images and storing the images
 - in the elastic file storage, which is how we will store data
 - going forward.

3 By giving each data source a fixed section of the overall video
 ↳ frame, we can display all the data in the final video by
 ↳ copying a given data source's image into its designated
 ↳ region. These regions would be hard coded.

4 For all tasks, reach out if anyone needs any help

5 Tasks for this week:

6 - someone should write a program that will assemble all the
 ↳ separate images (data sources) into one frame. Expect a
 ↳ progress update on tuesday on whether or not the task will
 ↳ need more time. Assigned to Ted and Ivan

7 - ongoing preparation of the AWS environment, especially so that
 ↳ it's completely ready for other parts of the project.
 ↳ Specifically, spinning up a docker container on a regular
 ↳ timer. Make sure to have the E2 and elastic storage set up.

8 - Garret will help Corbin with ongoing web scraping

9 - Edward will design the layout of all individual data sources,
 ↳ i.e. how big each box should be, where it should go, as
 ↳ concretely as possible.

5.5 Meeting notes 11/17

1 Members Absent: None

2 Next meeting: from 3:20 to 4:15 on Thursday, 11/19

3

4 Garrett and Corbin have finished with the web scraper, but some
 ↳ data (i.e. supernovae) are stored in plain text right now.
 Python scripts and selenium have been added to the git repo

5 TODO store text data as images, using the ffmpeg command (Ted
 6 ↳ worked on this). This should be relatively easy

7 Ivan and Ted completed their task of image stitching.

8 Their program requires input images to be .png using rgba pixel
 ↳ format.

9 `ffmpeg -i <filename.png>`

10 The above command can be used to check the pixel format

11 Edward is continuing work on finalizing the layout, including
 ↳ specific pixel measurements and placement

12 Marley has Cloud9 setup, as well as the elastic file storage, but
 ↳ there was some difficulty linking them together

13 The professor has said that he will make a tutorial on this
 ↳ process as many students are struggling with it, but no
 ↳ materials have been released as of yet

14

15 Tasks for next time

16 Corbin, Ted and Edward - Convert text data to images

17 Garrett, Ivan and Edward - Make a sample video from dummy images
 ↳ using image layout parameters

18 Time scaling: make sure that data are displayed for an
 ↪ appropriate amount of time, and at the correct time
19 Edward - finish the layout measurements by thursday
20 Marley - continue work on AWS related issues
21
22 Questions
23 Can we have help and/or a tutorial on Linking EFS with EC2? There
 ↪ seems to be some difficulties with security specifically.
24 See announcement from today

5.6 Meeting notes 11/24

1 cron and/or lambdas?
2 Cron requires continuous running of a machine
3 Lambdas do not
4 Split this into its own issue on the issue tracker
5 We would like to use lambda statements
6 AWS and EFS related tasks are complete
7 Marley posted instructions for working with EFS in the discord
 ↪ server: make sure to read and follow
8 EFS structure and naming conventions - directory framesegments
 ↪ has 6 subdirectories
9 issfeed
10 isscoords
11 lightningmap
12 supernovae
13 twitch
14 weather
15 frames
16
17 Ted demonstrated completed work on the text-to-image task. Due to
 ↪ power outages, sample output was not available.
18 Some supernovae data points occur with identical timestamps,
 ↪ which conflicts with the requirement that files be named
 ↪ after the unix time at which the data point occurred. We
 ↪ discussed and corrected this by having the script overwrite
 ↪ data points with identical times, so that only the last read
 ↪ data point is actually saved persistently.
19 Corbin is still working on parts of the text-to-image task.
20
21 When and how often should we start collecting data?
22 Supernovae: once per day
23 Other data sources: about once every 30 minutes
24 Start: ASAP
25 Where should docker files be stored?
26

27 Tasks for next week:
28 Scheduled automatic data collection via lambda statements -
 ↪ Marley and Ted
29 Ask Prof. Jensen about meeting next week - Edward
30 UML class and sequence diagrams - Edward
31 Current framebuilder program is hardcoded, and needs to be
 ↪ generalized to take arbitrary time ranges - Ivan and Garrett
32 Ongoing work on
33 An output video
34
35 Next meeting 1 PM sunday - let's keep it short please!
 ↪ (absolutely less than an hour)
36 Please report if you finish anything early
37 Review submission requirements

5.7 Meeting notes 11/29

1 Members absent: Ted
2
3 We commiserate over a lack of replies from Prof. Jensen...
4
5 Marley:
6 We really need prof. Help on this.
7 EC2 doesn't work on a schedule, so we use a cron job. Docker
 ↪ commands are run every 2 hours. It seems the amazon account
 ↪ logs us out at unpredictable times.
8 Automation of AWS stuff would require a new account without
 ↪ restrictions and some code refactoring
9 Lambda functions do work, but account restrictions prevent us
 ↪ from fully implementing it
10 Which commands need to go into the cron job file (to be run every
 ↪ 2 hours)?
11 Supernovae can be run once every 2 hours
12 Other sources can be run once every 30 minutes
13 For now, let's run data collection on a local machine using cron
 ↪ so we have something to show.
14
15 Garret and Ivan:
16 Generalizing the framebuilder software is going well and should
 ↪ be done tonight or tomorrow.
17
18 Edward: a bit behind on task. professor hasn't posted any
 ↪ instructions for assignment submission and documentation,
 ↪ which makes completing that requirement ambiguous at the
 ↪ moment.
19

20 Next meeting: tuesday at 2

5.8 Meeting notes 12/1

1 Garrett and Ivan are nearly done, scales correctly with time and
→ grabs images. Just needs command line arguments to specify a
→ time window.

2 Turn in date: let's set an appointment for 1 PM on the 10th

3 Edward has been busy and hasn't worked more on documentation yet,
→ but should be good moving forward.

4 Marley will send Edward a paragraph explaining troubleshooting
→ AWS problems to include in the documentation

5 Continued editing of the cron tab file

6 Text conversion python scripts should output to /dataStorage

7 We discussed some of the details of when docker containers should
→ run and how often.

8 Supernovae makes multiple files, so maybe run it only once every
→ two hours.

9 What Ted needs is a list of shell commands for each docker
→ container and how frequently each of them should be run.
→ Garrett can provide this.

10 Once the .sh files are done (Garrett and Ted), Marley can
→ schedule them with cron tab.

11 Goal: start data collection by tomorrow night

12 What path did we decide to mount EFS to?

13 Marley posted a screenshot in chat

14 Edward submitted times to meet with the professor

15 Edward posted on piazza asking for clarification on the documents
→ to be submitted

16 Next meeting: 2PM thursday

6 Appendix B: Layout design

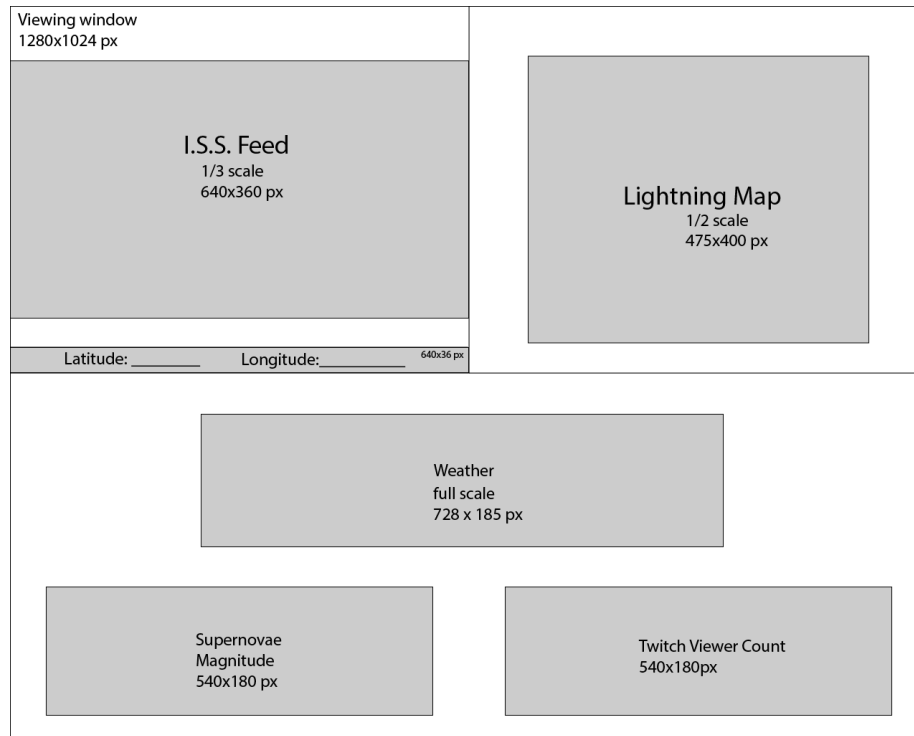


Figure 7: A finalized version of the layout from 11/10